


6.004 Tutorial Problems

L04 – Procedures and Stacks II

Symbolic name	Registers	Description	Saver
a0 to a7	x10 to x17	Function arguments	Caller
a0 and a1	x10 and x11	Function return values	Caller
ra	x1	Return address	Caller
t0 to t6	x5-7, x28-31	Temporaries	Caller
s0 to s11	x8-9, x18-27	Saved registers	Callee
sp	x2	Stack pointer	Callee
gp	x3	Global pointer	---
tp	x4	Thread pointer	---

RISC-V Calling Conventions:

- Caller places arguments in registers a0–a7
- Caller transfers control to callee using jal (jump-and-link) to capture the return address in register ra. The following three instructions are equivalent (pc stands for program counter, the memory address of the current/next instruction):
 - jal ra, label: R[ra] <= pc + 4; pc <= label
 - jal label (pseudoinstruction for the above)
 - call label (pseudoinstruction for the above)
- Callee runs, and places results in registers a0 and a1
- Callee transfers control to caller using jr (jump-register) instruction. The following instructions are equivalent:
 - jalr x0, 0(ra): pc <= R[ra]
 - jr ra (pseudoinstruction for the above)
 - ret (pseudoinstruction for the above)



Push register **x_i** onto stack
`addi sp, sp, -4`
`sw xi, 0(sp)`

Pop value at top of stack into register **x_i**
`lw xi, 0(sp)`
`addi sp, sp, 4`


Assume 0(sp) holds valid data.

Stack discipline: can put anything on the stack, but leave stack the way you found it

Always save **s** registers before using them
 Save **a** and **t** registers if you will need their value after procedure call returns.
 Always save **ra** if making nested procedure calls.

RISC-V Stack

- Stack is in memory → need a register to point to it
 - In RISC-V, stack pointer `sp` is `x2`
- Stack grows down from higher to lower addresses
 - Push decreases `sp`
 - Pop increases `sp`
- `sp` points to top of stack (last pushed element)
- Discipline: Can use stack *at any time*, but leave it as you found it!



February 12, 2020

MIT 6.004 Spring 2020

L03-19

Using the stack

Sample entry sequence

```
addi sp, sp, -8
sw ra, 0(sp)
sw a0, 4(sp)
```

Corresponding Exit sequence

```
lw ra, 0(sp)
lw a0, 4(sp)
addi sp, sp, 8
```

February 12, 2020

MIT 6.004 Spring 2020

L03-20

Note: A small subset of essential problems are marked with a red star (★). We especially encourage you to try these out before recitation.

Problem 1.

Write assembly program that computes square of the sum of two numbers (i.e. $\text{squareSum}(x,y) = (x + y)^2$) and follows RISC-V calling convention. Note that in your assembly code you have to call assembly procedures for **mult** and **sum**. They are not provided to you, but they are fully functional and obey the calling convention.

```
/* compute square sum of args */
unsigned squareSum(unsigned x, unsigned y) {
    unsigned z = sum(x, y);
    return mult(z, z);
}
// start of assembly code
```

Handwritten annotations: a_0 points to `sum(x, y)`, a_1 points to `mult(z, z)`. a_0 and a_1 are also written below the code.

① Figure registers to store in mem.

```
addi    sp, sp, -4
sw      ra, 0(sp)
jal     sum
mv      a1, a0
jal     mult
lw      ra, 0(sp)
addi    sp, sp, 4
ret
```

Handwritten annotations: a_0 is underlined. $-ra$ is underlined.

★ $i \log_2 \text{stack} \uparrow 2$

$\Delta i \log_2 \text{stack} \uparrow 2$

$0, i \log_2 \text{stack} \uparrow 2$
 \vdots
 $\text{stack} \downarrow 2$

Problem 2. ★

The following C program computes the log base 2 of its argument. The assembly code for the procedure is shown on the right, along with a stack trace showing the execution of `ilog2(10)`. The execution has been halted just as it's about to execute the instruction labeled `rtn`. The SP label on the stack shows where the SP is pointing to when execution halted.

```
/* compute log base 2 of arg */
int ilog2(unsigned x) {
    unsigned y;
    if (x == 0) return 0;
    else {
        /* shift x right by 1 bit */
        y = x >> 1;
        return ilog2(y) + 1;
    }
}
```

```
ilog2: beqz a0, rtn
        addi sp, sp, -8
        sw s0, 4(sp)
        sw ra, 0(sp)
        srli s0, a0, 1
        mv a0, s0
        jal ra, ilog2
        addi a0, a0, 1
        lw ra, 0(sp)
        lw s0, 4(sp)
        addi sp, sp, 8
        rtn: jr ra
```

(A) Please fill in the values for the two blank locations in the stack trace shown on the right. Please express the values in hex.

Fill in values (in hex!) for 2 blank locations

(B) What are the values in `a0`, `s0`, `sp`, and `pc` at the time execution was halted? Please express the values in hex or write "CAN'T TELL".


Value in `a0`: 0x 2 in `s0`: 0x 2

Value in `sp`: 0x Can't tell in `pc`: 0x 250

(C) What was the address of the original `ilog2(10)` function call?

Original `ilog2(10)` address: 0x 1104

0x0	
0x93	
0x240	(ra)
0x1	s0
0x240	(ra)
0x2	s0
SP	
0x240	(ra)
0x5	s0
0x1108	ra
0x37	s0



0b 00000001 0b 00000001
00000000 00000001

Problem 3. ★

You are given an incomplete listing of a C program (shown below) and its translation to RISC-V assembly code (shown on the right):

```
int fn(int x) {
    int lowbit = x & 1;
    int rest = x >> 1;
    if (x == 0) return 0;
    else return ???;
}
```

anding
 $s0 = a0 \& 1$
 $lowbit = x \& 1$

rest = s1

```
fn: addi sp, sp, -12
     sw s0, 0(sp)
     sw s1, 4(sp)
     sw ra, 8(sp)
     andi s0, a0, 1
     srai s1, a0, 1
     yy: beqz a0, rtn
        mv a0, s1
        jal ra, fn
        add a0, a0, s0
```

a0

(A) What is the missing C source corresponding to ??? in the above program?

C source code:

$fn(rest) + lowbit$

```
rtn: lw s0, 0(sp)
     lw s1, 4(sp)
     lw ra, 8(sp)
     addi sp, sp, 12
     jr ra
```

rest
 $return\ fn(\text{rest}) + lowbit$

fnct 1

base

fnct 1

0b1001100

+0 0b100110

+0 0b010011

+1

0b001001

+1

0b00100

+0

0b10

+0

0b1

(3) $(+1)$ 660000

The diagram shows a memory stack with the following addresses and values:

Address	Value
0x1D0	0
0x1C0	0x1
0x1B0	0x11
0x1A0	0x4C
0x190	0x23
0x180	0x4C
0x170	0x3
0x160	0x22
0x150	0xC4

Annotations and instructions:

- SP** (Stack Pointer) is indicated by a red bracket on the left, pointing to the 0x1D0 address.
- SP →** is written next to the 0x1B0 address.
- Instructions** (written in green) are: *sub \$0, \$1, \$ra* (pointing to 0x1A0), *sub \$0, \$1, \$ra* (pointing to 0x180), and *sub \$0, \$1, \$ra* (pointing to 0x150).
- Registers** (written in green) are: *\$0*, *\$1*, and *\$ra*.
- Other annotations** include a green arrow pointing up from the 0x1A0 address, a green arrow pointing right from the 0x1B0 address to the value 0x11, and a green arrow pointing down from the 0x1A0 address to the value 0x4C.

Most recent argument (HEX): x= 0x11

Contents of 1D0 (HEX): Can't tell

Address of rtn (HEX): 0x50

First recursive call argument (HEX): $x = 0x23$

Address of original call (HEX): $0 \times C 0$

Original s1 contents (HEX): 0x22

Return value for original call (HEX): _____

6.004 Fall 2021 Worksheet

Problem 4. ★

The following C program implements a function $H(x,y)$ of two arguments, which returns an integer result. The assembly code for the procedure is shown on the right.

```
int H(int x, int y) {
    int a = x - y;
    if (a < 0) return x;
    else return ???;
}
```

```
H:  sub t0, a0, a1
    bltz t0, rtn
    addi sp, sp, -4
    sw ra, 0(sp)
    mv a0, t0
    jal H
    lw ra, 0(sp)
    addi sp, sp, 4
rtn: jr ra
```

The execution of the procedure call $H(0x68, 0x20)$ has been suspended just as the processor is about to execute the instruction labeled “rtn.” **during one of the recursive calls to H.** A *partial* trace of the stack at the time execution was suspended is shown to the right below.

(A) Examining the assembly language for H, what is the appropriate C code for ??? in the C representation for H?

C code for ???:

$H(a_0, a_1)$

(B) Please fill in the values for the blank locations in the stack dump shown on the right. Express the values in hex or write “---” if value can’t be determined. For all following questions, suppose that during the initial (non-recursive) call to H, sp pointed to the memory location containing 0x0010.

Fill in the blank locations with values (in hex!) or “---”

0x007c
sp →
0x001c
0x0010
0x00e0
0x00ec

(C) Determine the specified values at the time execution was suspended. Please express each value in hex or write “CAN’T TELL” if the value cannot be determined.

Value in a0 or “CANT TELL”: 0x_____

Value in a1 or “CANT TELL”: 0x_____

Value in ra or “CANT TELL”: 0x_____

Value in sp or “CANT TELL”: 0x_____

Address of the initial call instruction to H: 0x_____

From past quizzes:

Problem 4. Stack Detective (16 points)

Below is the Python code for a recursive implementation of binary search, which finds the index at which an element should be inserted into a sorted array to ensure that the array is still sorted after the insertion. To the right is a not so elegant, but valid, implementation of the function using RISC-V assembly.

```
/* find where to insert element in arr */
unsigned binary_search(int[] arr,
                      unsigned start,
                      unsigned end,
                      int element){
    if (start == end){
        return end;
    }
    mid = (start + end) / 2;
    if (element < arr[mid]){
        end = mid;
    } else {
        start = mid + 1;
    }
    return binary_search(arr, start, end,
                        element);
}
```

```
binary_search: addi sp, sp, -8
               sw ra, 4(sp)
               sw s0, 0(sp)
               mv s0, a2
               beq a1, a2, done
               add t0, a1, a2
               srli t0, t0, 1
               slli t1, t0, 2
               add t1, a0, t1
               lw t1, 0(t1)
if:            bge _____
               mv a2, t0
               j recurse
else:          addi t0, t0, 1
               mv a1, t0
recurse:       call binary_search
               mv s0, a0
done:          mv a0, s0
               lw s0, 0(sp)
               lw ra, 4(sp)
L1:           addi sp, sp, 8
               ret
```

- (A) (2 points) What should be in the blank on the line labeled **if** to make the assembly implementation match the Python code?

if: bge _____

- (B) (2 points) How many words will be written to the stack before the program makes each recursive call to the function **binary_search**?

Number of words pushed onto stack before each recursive call? _____

The program's initial call to function **binary_search** occurs outside of the function definition via the instruction '**call binary_search**'. The program is interrupted *during a recursive call* to **binary_search**, *just prior* to the execution of '**addi sp, sp, 8**' at label **L1**. The diagram on the right shows the contents of a region of memory. All addresses and data values are shown in hex. **The current value in the SP register is 0xEB0 and points to the location shown in the diagram.**

- (C) (4 points) What were the values of arguments **arr** and **end** at the beginning of *the initial call* to **binary_search**? Write CAN'T TELL if you cannot tell the value of an argument from the stack provided.

Arguments at beginning of this call : **arr** = 0x_____

end = 0x_____

- (D) (4 points) What are the values in the following registers right when the execution of **binary_search** is interrupted? Write CAN'T TELL if you cannot tell.

Current value of **s0**: 0x_____

Current value of **ra**: 0x_____

Memory Contents

Address	Data
0xEA4	0x0
0xEA8	0x5
0xEAC	0xC4
0xEB0	0x6
0xEB4	0xC4
0xEB8	0x6
0xEBC	0xC4
0xEC0	0xA
0xEC4	0xC4
0xEC8	0x3E
0xECC	0xCA4
0xED0	0xCED

SP→

- (E) (2 points) What is the hex address of the '**call binary_search**' instruction that made the *initial call* to **binary_search**?

Address of instruction that made initial call to **binary_search**: 0x_____

- (F) (2 points) What is the hex address of the **ret** instruction?

Address of **ret** instruction: 0x_____